



# 嵌入式系統總整與實作

曾煜棋、吳昆儒

**National Yang Ming Chiao Tung University**



日期	主題
2/21	0. 課程介紹
2/28	梅竹賽!!
3/7	1. 嵌入式開發板 - 樹莓派介紹與設定 (headless)
3/14	2. 連接感測器 (GPIO, I2C)
3/21	3. 整合感測資訊 (IMU)
3/28	4. 整合音訊資料 (麥克風, 語音識別)
4/4	清明節放假
4/11	5. 整合視覺資料 (攝影機, 影像辨識)
4/18	期中考Midterm, Project分組
4/25	6. 嵌入式模型 (mediapipe, video, audio, text)
5/2	7. 喚醒詞原理 (by 台灣樹莓派)
5/9	<b>Final Project – Proposal 分組報告 (online?)</b>
5/16	8. 樹莓派核心編譯 (Cross compile, Kernel)
5/23	9. 嵌入式套件編譯
5/30	端午節放假
6/6	Final project準備周, Q&A (學期考試周)
<b>6/13</b>	<b>Final Project demonstration</b>

期中考周  
(4/7-4/11)

期末考周  
(6/2-6/6)



# Last week

- 整合視覺資料 (攝影機, 影像辨識)
  - Part1: 安裝相機模組
  - Part2: 人臉偵測, 輪廓偵測
  - Part3: 物件辨識 (tensorflow, YOLO)
- Install dependency
  - `sudo apt install python3-opencv`
  - `pip install imutils`
  - `pip install dlib`



# This week

- 嵌入式模型 (mediapipe)
  - Part1: Mediapipe intro
  - Part2: Text classification
  - Part3: Audio classification
  - Part4: Image applications (Pose estimation)



Input:

“Great movie with a classic plot. I will recommend this to everyone.”

Output:




Bird Call





# MediaPipe Solutions

- MediaPipe Solutions offers a suite of AI/ML libraries and tools
- Easily and quickly integrates into various applications
- Supports multiple development platforms





# Available solutions

- LLM Inference API
- Object detection
- Image classification
- Image segmentation
- Interactive segmentation
- Hand landmark detection
- Gesture recognition
- Image embedding
- Face detection
- Face landmark detection
- Face stylization
- Pose landmark detection
- Image generation
- Text classification
- Text embedding
- Language detector
- Audio classification



# MediaPipe

- **Text Classification**

- Enter text in the command line and classify it as either positive or negative with a provided confidence score

- **Audio classifier**

- perform real-time audio classification using audio streamed from the microphone.

- **Image Application**

- Including Face detection, Object detection, Pose landmark detection ... etc



# 注意OS版本

- Note: Before you begin, you need to set up your Raspberry Pi with Raspberry 64-bit Pi OS.

## 1. 安裝OS (Raspbian)

### 步驟1：下載映像檔

- 官方下載網頁
  - 官方連結 (Raspberry Pi OS 64-bit (Legacy) with desktop)
    - [https://downloads.raspberrypi.com/raspios\\_oldstable\\_armhf/images/raspios\\_oldstable\\_armhf-2024-10-28/2024-10-22-raspios-bullseye-armhf.img.xz](https://downloads.raspberrypi.com/raspios_oldstable_armhf/images/raspios_oldstable_armhf-2024-10-28/2024-10-22-raspios-bullseye-armhf.img.xz)

#### Raspberry Pi OS (Legacy, 64-bit)

Compatible with:

3B 3B+ 3A+ 4B  
400 CM3 CM3+  
CM4 CM4S Zero 2 W

#### Raspberry Pi OS (Legacy) with desktop

Release date: October 22nd 2024

System: 64-bit

Kernel version: 6.1

Debian version: 11 (bullseye)

Size: 868MB

[Show SHA256 file integrity hash:](#)

[Release notes](#)

[Download](#)

[Download torrent](#)

[Archive](#)





# 需重新準備OS

- 下載最新版64bit OS
  - [https://downloads.raspberrypi.com/raspios\\_arm64/images/raspios\\_arm64-2024-11-19/2024-11-19-raspios-bookworm-arm64.img.xz](https://downloads.raspberrypi.com/raspios_arm64/images/raspios_arm64-2024-11-19/2024-11-19-raspios-bookworm-arm64.img.xz)
- 設定headless參數 (cmdline.txt, config.txt)
- 建立使用者 (userconf.txt)
- 開啟SSH服務 (filename with SSH)

本周五會提供新的SD卡給大家用  
(已安裝OS與mediapipe + headless設定)



# Python環境注意事項

- **Starting in Raspberry Pi OS Bookworm, packages installed via pip must be installed into a Python virtual environment (venv).**
- A virtual environment is a container where you can safely install third-party modules so they won't interfere with your system Python.



# 直接安裝會有錯誤訊息

```
pi@raspberrypi:~ $ pip install mediapipe  
error: externally-managed-environment
```

× This environment is externally managed

↳ To install Python packages system-wide, try apt install python3-xyz, where xyz is the package you are trying to install.

If you wish to install a non-Debian-packaged Python package, create a virtual environment using python3 -m venv path/to/venv. Then use path/to/venv/bin/python and path/to/venv/bin/pip. Make sure you have python3-full installed.

For more information visit <http://rptl.io/venv>

note: If you believe this is a mistake, please contact your Python installation or OS distribution provider. You can override this, at the risk of breaking your Python installation or OS, by passing --break-system-packages.


hint: See PEP 668 for the detailed specification.




# Virtualenv

- virtualenv is a tool to create isolated Python environments


Virtual Environment 01



Virtual Environment 02



Virtual Environment 03



simple**LEARN**



# Mediapipe venv

- `sudo apt-get update`
- `sudo apt-get install -y python3-dev python3-pip`
- `python -m venv --system-site-packages mdp`
- **`source ~/mdp/bin/activate`      // 進入虛擬環境**
- `pip install matplotlib opencv-python imutils`
- `pip install mediapipe`
- `git clone https://github.com/google-ai-edge/mediapipe-samples`
- .....
- **`deactivate`      // 離開虛擬環境**



**Input:**

“Great movie with a classic plot. I will recommend this to everyone.”

**Output:**



## Text Classification

[https://github.com/google-ai-edge/mediapipe-samples/tree/main/examples/text\\_classification/raspberry\\_pi](https://github.com/google-ai-edge/mediapipe-samples/tree/main/examples/text_classification/raspberry_pi)



# Text Classification

- This model will accept text entered in the command line and classify it as either **positive** or **negative** with a provided confidence score.
- The models were trained on the SST-2 (Stanford Sentiment Treebank) dataset, which consists of movie reviews labeled as either positive or negative.
- The supported classification models include Average Word-Embedding and MobileBERT.



# Classification models

- **BERT-classifier model**
  - Based on MobileBERT; offers high accuracy
  - Includes metadata for BERT tokenization
  - Recommended for scenarios needing higher accuracy
- **Average word embedding model**
  - Uses average word embedding; smaller and faster
  - Lower prediction accuracy, but faster to retrain
  - Includes metadata for regex tokenization
  - Suitable for resource-limited or efficiency-focused scenarios

MobileBERT: a Compact Task-Agnostic BERT for Resource-Limited Devices  
<https://arxiv.org/abs/2004.02984>

Model Name	CPU Latency
Average word embedding	0.14ms
BERT-classifier	57.68ms





# Run the example

- `git clone https://github.com/google-ai-edge/mediapipe-samples`
- `cd mediapipe-samples/examples/text_classification/raspberry_pi`
- `sh setup.sh`
- `python3 classify.py --inputText "Your text goes here"`

```
(mdp) pi@raspberrypi:~/mediapipe-samples/examples/text_classification/raspberry_pi $ python3 classify.py --inputText "Today is a wonderful day to build something people love"
Error in cpuinfo: prctl(PR_SVE_GET_VL) failed
INFO: Created TensorFlow Lite XNNPACK delegate for CPU.
WARNING: All log messages before absl::InitializeLog() is called are written to STDERR
W0000 00:00:1745167556.943222 5747 inference_feedback_manager.cc:114] Feedback manager requires a model with a single signature inference. Disabling support for feedback tensors.
positive (1.00)
```



# Discussion 1

- The sample code use BERT-classifier model by default.
- How to use **average word embedding model** to classify the text?



# Quiz 1

- Integrate Microphone Input, Speech-to-Text, and Text Classification
- Capture real-time audio from a microphone, transcribe the speech to text, and then classify the resulting text using a text classification model
- You have to rebuild environment
  - pip install SpeechRecognition
  - pip install openai
- Hint: you can use python's **subprocess**

```
Please wait. Calibrating microphone...
Say something!
OpenAI Whisper API thinks you said: The quick fox jumps over a lazy dog.
classify.py stdout:
negative (0.64)
```



# Speech-to-Text example

```
1 import speech_recognition as sr
2 import os
3
4 #obtain audio from the microphone
5 r=sr.Recognizer()
6
7 with sr.Microphone() as source:
8     print("Please wait. Calibrating microphone...")
9     #listen for 1 seconds and create the ambient noise energy level
10    r.adjust_for_ambient_noise(source, duration=1)
11    print("Say something!")
12    audio=r.listen(source)
13
14 # recognize speech using Whisper API
15 OPENAI_API_KEY = " "
16 os.environ["OPENAI_API_KEY"] = OPENAI_API_KEY
17 try:
18     print(f"OpenAI Whisper API thinks you said {r.recognize_openai(audio)}")
19 except sr.RequestError as e:
20     print(f"Could not request results from OpenAI Whisper API; {e}")
21
```



## Audio classifier

[https://github.com/google-ai-edge/mediapipe-samples/tree/main/examples/audio\\_classifier/raspberry\\_pi](https://github.com/google-ai-edge/mediapipe-samples/tree/main/examples/audio_classifier/raspberry_pi)



# Audio classification

- This model **classifies audio clips** into a set of **defined categories**, such as **guitar music, a train whistle, or a bird's song**.
- The categories are defined during the training of the model.
- This task operates on audio data with a machine learning (ML) model as independent audio clips or a continuous stream and outputs a list of potential categories ranked by descending probability score.



# Classification models

- **Yamnet model**


- This audio event classifier trained on the AudioSet dataset to predict audio events defined in the AudioSet data.

- **AudioSet dataset**

- A large-scale dataset of manually annotated audio events

Audio Set: An ontology and human-labeled dataset for audio events

<https://ieeexplore.ieee.org/document/7952261>




Model Name	CPU Latency
YamNet	12.29ms



# Run the example

- git clone <https://github.com/google-ai-edge/mediapipe-samples>
- cd mediapipe/examples/audio\_classifier/raspberry\_pi
- sh setup.sh
- python3 classify.py

```
(mdp) pi@raspberrypi:~/mediapipe-samples/examples/audio_classifier/raspberry_pi $ python3 classify.py
Error in cpuinfo: prctl(PR_SVE_GET_VL) failed
INFO: Created TensorFlow Lite XNNPACK delegate for CPU.
WARNING: All log messages before absl::InitializeLog() is called are written to STDERR
W0000 00:00:1745205146.017366    6428 inference_feedback_manager.cc:114] Feedback manager requires a model with a single signature inference. Disabling support for feedback tensors.
W0000 00:00:1745205147.092885    6428 time_series_util.cc:52] Timestamp 1745205147.090000 not consistent with number of samples 15600 and initial timestamp 1745205146601000. Expected timestamp: 1745205147.576000 Timestamp difference: -0.486000 sample_rate: 16000.000000
W0000 00:00:1745205159.628734    6429 time_series_util.cc:52] Timestamp 1745205159.628000 not consistent with number of samples 327600 and initial timestamp 1745205146601000. Expected timestamp: 1745205167.076000 Timestamp difference: -7.448000 sample_rate: 16000.000000
```











# Discussion 2

- AudioSet consists of an expanding ontology of 632 audio event classes and a collection of 2,084,320 human-labeled 10-second sound clips drawn from YouTube videos.
- How can we find the **full list** of audio events?





## Image Application

[https://github.com/google-ai-edge/mediapipe-samples/tree/main/examples/face\\_detector/raspberry\\_pi](https://github.com/google-ai-edge/mediapipe-samples/tree/main/examples/face_detector/raspberry_pi)


[https://github.com/google-ai-edge/mediapipe-samples/tree/main/examples/object\\_detection/raspberry\\_pi](https://github.com/google-ai-edge/mediapipe-samples/tree/main/examples/object_detection/raspberry_pi)

[https://github.com/google-ai-edge/mediapipe-samples/tree/main/examples/pose\\_landmarker/raspberry\\_pi](https://github.com/google-ai-edge/mediapipe-samples/tree/main/examples/pose_landmarker/raspberry_pi)



# Face detection


- This model **detects faces** in an image or video.
- You can use this task to locate faces and facial features within a frame.
- The task outputs **face locations**, along with the following **facial key points**: left eye, right eye, nose tip, mouth, left eye tragion, and right eye tragion.





# Models

- The models are variants of **BlazeFace**, a lightweight and accurate face detector optimized for mobile GPU inference.
- BlazeFace models are suitable for applications like 3D facial keypoint estimation, expression classification, and face region segmentation.



BlazeFace: Sub-millisecond Neural Face Detection on Mobile GPUs  
<https://arxiv.org/pdf/1907.05047>

Figure 3. Pipeline example (best viewed in color).  
Red: BlazeFace output. Green: Task-specific model output.




# Comparisons

- **Haar Cascade**
  - Uses hand-crafted Haar features and combines many weak classifiers with AdaBoost.
  - Sensitive to lighting and pose changes, leading to more false detections.
- **HOG (Histogram of Oriented Gradients)**
  - Extracts edge orientation features and classifies using an SVM.
  - Less robust to pose, lighting, and background variations; not ideal for real-time.
- **BlazeFace**
  - Uses a lightweight CNN for fast and accurate real-time face detection.
  - Well-suited for mobile devices and can handle faces from multiple angles.



# Object detection

- This model detects the presence and location of multiple classes of objects within images or videos.
- Each detection result represents an object that appears within the image or video.





# Models

- **EfficientDet-Lite0 (Recommended)**
  - Balanced in accuracy and speed, suitable for many lightweight use cases.
  - Uses a 320x320 input with EfficientNet-Lite0 backbone; supports int8/float16/float32.
- **EfficientDet-Lite2**
  - Achieves higher accuracy than Lite0 but is slower and uses more memory.
  - Uses a 448x448 input with EfficientNet-Lite2 backbone; suitable when accuracy is the top priority.
- **SSD MobileNetV2**
  - Fastest and lightest among the three models but less accurate.
  - Uses a 256x256 input with MobileNetV2 backbone; ideal when speed and small size are critical.

EfficientDet: Scalable and Efficient Object Detection


[https://openaccess.thecvf.com/content\\_CVPR\\_2020/html/Tan\\_EfficientDet\\_Scalable\\_and\\_Efficient\\_Object\\_Detection\\_CVPR\\_2020\\_paper.html](https://openaccess.thecvf.com/content_CVPR_2020/html/Tan_EfficientDet_Scalable_and_Efficient_Object_Detection_CVPR_2020_paper.html)





# Pose landmark detection

- The MediaPipe Pose Landmarker task lets you **detect landmarks of human bodies** in an image or video.
- You can use this task to identify **key body locations**, analyze posture, and categorize movements.
- The task outputs body pose landmarks in image coordinates and in 3-dimensional world coordinates.







# Models

- The Pose Landmarker uses a series of models to predict pose landmarks.
  - **Pose detection model:** detect the presence of bodies with a few key pose landmarks.
  - **Pose landmarker model:** add a complete mapping of the pose. The model outputs an estimate of 33 3-dimensional pose landmarks.





Figure 1. Inference pipeline. See text.



# Pose landmarker model



- |                       |                       |
|-----------------------|-----------------------|
| 0 - nose              | 16 - right wrist      |
| 1 - left eye (inner)  | 17 - left pinky       |
| 2 - left eye          | 18 - right pinky      |
| 3 - left eye (outer)  | 19 - left index       |
| 4 - right eye (inner) | 20 - right index      |
| 5 - right eye         | 21 - left thumb       |
| 6 - right eye (outer) | 22 - right thumb      |
| 7 - left ear          | 23 - left hip         |
| 8 - right ear         | 24 - right hip        |
| 9 - mouth (left)      | 25 - left knee        |
| 10 - mouth (right)    | 26 - right knee       |
| 11 - left shoulder    | 27 - left ankle       |
| 12 - right shoulder   | 28 - right ankle      |
| 13 - left elbow       | 29 - left heel        |
| 14 - right elbow      | 30 - right heel       |
| 15 - left wrist       | 31 - left foot index  |
|                       | 32 - right foot index |




# Run the example

- `git clone https://github.com/google-ai-edge/mediapipe-samples`
- `cd mediapipe/examples/pose_landmarker/raspberry_pi`
- `sh setup.sh`
- `python3 detect.py`
- Note: This example uses MediaPipe with Python on a Raspberry Pi to perform real-time pose landmarks detection using images streamed from the camera.



# Run sample code

- mp\_image.py: read a image to calculate keypoints
- Usage: `python mp_image.py your_image`



```
(mdp) pi@raspberrypi:~/mediapipe-samples/examples/pose_landmarker/raspberry_pi $ python mp_image.py warrior-pose.jpg
Downloading model to /home/pi/mdp/lib/python3.11/site-packages/mediapipe/modules/pose_landmarker/pose_landmark_heavy.tflite
Error in cpuinfo: prctl(PR_SVE_GET_VL) failed
INFO: Created TensorFlow Lite XNNPACK delegate for CPU.
WARNING: All log messages before absl::InitializeLog() is called are written to STDERR
W0000 00:00:1745220846.891218    7364 inference_feedback_manager.cc:114] Feedback manager requires a model with a single signature inference. Disabling support for feedback tensors.
W0000 00:00:1745220847.230398    7367 inference_feedback_manager.cc:114] Feedback manager requires a model with a single signature inference. Disabling support for feedback tensors.
W0000 00:00:1745220847.771356    7367 landmark_projection_calculator.cc:186] Using NORM_RECT without IMAGE_DIMENSIONS is only supported for the square ROI. Provide IMAGE_DIMENSIONS or use PROJECTION_MATRIX.
=====
0 NOSE (534, 141, -107.0648655295372, 0.9999992847442627)
=====
1 LEFT_EYE_INNER (528, 131, -88.58142793178558, 0.9999954700469971)
=====
2 LEFT_EYE (527, 131, -88.80846947431564, 0.999996060934448)
=====
3 LEFT_EYE_OUTER (526, 131, -88.7698233127594, 0.9999959468841553)
=====
4 RIGHT_EYE_INNER (526, 131, -112.45574802160263, 0.999997615814209)
=====
5 RIGHT_EYE (523, 131, -112.72476613521576, 0.9999983310699463)
=====
6 RIGHT_EYE_OUTER (519, 132, -112.77415603399277, 0.9999980926513672)
=====
7 LEFT_EAR (508, 137, -7.074025459587574, 0.9999982118606567)
```



## 載入套件與初始化模型

```
#!/usr/bin/python3
import cv2
import sys
import time
import numpy as np
import mediapipe as mp

mp_drawing = mp.solutions.drawing_utils
mp_pose = mp.solutions.pose
pose = mp_pose.Pose(static_image_mode=True, min_detection_confidence=0.3, model_complexity=2)
```

```
try:
    image = cv2.imread(sys.argv[1])
except Exception as e:
    print(e)
    image = cv2.imread("images/warrior-pose.jpg")
```

## 讀取圖片

```
output_image = image.copy()
imageRGB = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
results = pose.process(imageRGB)
height, width, _ = image.shape
landmarks = []
```

## 姿態辨識

```
# Check if any landmarks are detected.
if results.pose_landmarks:
    # Draw Pose landmarks on the output image.
    mp_drawing.draw_landmarks(image=output_image, landmark_list=results.pose_landmarks,
                              connections=mp_pose.POSE_CONNECTIONS)
    # Iterate over the detected landmarks.
    for landmark in results.pose_landmarks.landmark:
        # Append the landmark into the list.
        landmarks.append((int(landmark.x * width), int(landmark.y * height),
                          (landmark.z * width), landmark.visibility))

for idx, landmark in enumerate(landmarks):
    print("=====")
    print(idx, mp_pose.PoseLandmark(idx).name, landmark)
```

## 輸出座標(x,y,z)


```
cv2.imshow("preview", output_image)
cv2.waitKey(0)
cv2.destroyAllWindows()
```

## 顯示結果



# 姿態識別應用


- APP: 健身比對姿勢, 復健治療, 運動強化追蹤



<https://www.luxurywatcher.com/zh-Hant/article/29928>



# 姿態分類基礎




<https://developers.google.com/ml-kit/vision/pose-detection/classifying-poses>





# 姿態分類

- 計算多個關鍵點間的角度 + 條件判斷
- 以Warrior II Pose為例
  - $195 > \text{left\_elbow\_angle} > 165$
  - $110 > \text{left\_shoulder\_angle} > 80$
  - $195 > \text{left\_knee\_angle} > 165$
  - $120 > \text{right\_knee\_angle} > 90$








# 判斷姿勢sample code

- mp\_pose\_compare.py: 判斷姿勢
- Usage: python mp\_pose\_compare.py warrior-pose-ii.jpg

```
=====
left_elbow_angle: 181
right_elbow_angle: 186
# 1. Both arms are straight
=====
left_shoulder_angle: 90
right_shoulder_angle: 98
# 2. Shoulders are at the required angle
=====
left_knee_angle: 175
right_knee_angle: 109
# 3. One leg is straight
=====
left_knee_angle: 175
right_knee_angle: 109
# 4. The other leg is bent at the required angle
=====
```






# 計算三個點之間的夾角

```
def calculateAngle(landmark1, landmark2, landmark3):  
    x1, y1, _ = landmark1  
    x2, y2, _ = landmark2  
    x3, y3, _ = landmark3  
  
    angle = math.degrees(math.atan2(y3 - y2, x3 - x2) - math.atan2(y1 - y2, x1 - x2))  
  
    if angle < 0:  
        angle += 360  
  
    angle = min(angle, 360 - angle)  
  
    return angle
```

夾角等於兩個向量的arctan差值



# 挑選關節點



```
#-----  
# Calculate the required angles.  
#-----  
  
try:  
    # Get the angle between the left shoulder, elbow and wrist points.  
    left_elbow_angle = calculateAngle(landmarks[mp_pose.PoseLandmark.LEFT_SHOULDER.value],  
                                      landmarks[mp_pose.PoseLandmark.LEFT_ELBOW.value],  
                                      landmarks[mp_pose.PoseLandmark.LEFT_WRIST.value])  
  
    # Get the angle between the right shoulder, elbow and wrist points.  
    right_elbow_angle = calculateAngle(landmarks[mp_pose.PoseLandmark.RIGHT_SHOULDER.value],  
                                       landmarks[mp_pose.PoseLandmark.RIGHT_ELBOW.value],  
                                       landmarks[mp_pose.PoseLandmark.RIGHT_WRIST.value])  
  
    # Get the angle between the left elbow, shoulder and hip points.  
    left_shoulder_angle = calculateAngle(landmarks[mp_pose.PoseLandmark.LEFT_ELBOW.value],  
                                         landmarks[mp_pose.PoseLandmark.LEFT_SHOULDER.value],  
                                         landmarks[mp_pose.PoseLandmark.LEFT_HIP.value])  
  
    # Get the angle between the right hip, shoulder and elbow points.  
    right_shoulder_angle = calculateAngle(landmarks[mp_pose.PoseLandmark.RIGHT_HIP.value],  
                                          landmarks[mp_pose.PoseLandmark.RIGHT_SHOULDER.value],  
                                          landmarks[mp_pose.PoseLandmark.RIGHT_ELBOW.value])  
  
    # Get the angle between the left hip, knee and ankle points.  
    left_knee_angle = calculateAngle(landmarks[mp_pose.PoseLandmark.LEFT_HIP.value],  
                                     landmarks[mp_pose.PoseLandmark.LEFT_KNEE.value],  
                                     landmarks[mp_pose.PoseLandmark.LEFT_ANKLE.value])  
  
    # Get the angle between the right hip, knee and ankle points  
    right_knee_angle = calculateAngle(landmarks[mp_pose.PoseLandmark.RIGHT_HIP.value],  
                                      landmarks[mp_pose.PoseLandmark.RIGHT_KNEE.value],  
                                      landmarks[mp_pose.PoseLandmark.RIGHT_ANKLE.value])
```



# 檢查姿勢



```
#-----  
# Check if it is the warrior II pose  
#-----  
  
print("=====  
# 1. Check if the both arms are straight.  
print("left_elbow_angle:", int(left_elbow_angle))  
print("right_elbow_angle:", int(right_elbow_angle))  
if 165 < left_elbow_angle < 195 and 165 < right_elbow_angle < 195:  
    print("# 1. Both arms are straight")  
    print("=====  
  
print("left_shoulder_angle:", int(left_shoulder_angle))  
print("right_shoulder_angle:", int(right_shoulder_angle))  
# 2. Check if shoulders are at the required angle.  
if 80 < left_shoulder_angle < 110 and 80 < right_shoulder_angle < 110:  
    print("# 2. Shoulders are at the required angle")  
    print("=====  
  
print("left_knee_angle:", int(left_knee_angle))  
# 3. Check if one leg is straight.  
if 165 < left_knee_angle < 195:  
    print("# 3. One leg is straight")  
    print("=====  
  
# 4. Check if the other leg is bended at the required angle.  
print("right_knee_angle", int(right_knee_angle))  
if 90 < right_knee_angle < 120:  
    print("# 4. The other leg is bended at the required angle")  
    print("=====  
  
# Specify the label of the pose that is Warrior II pose.  
label = 'Warrior II Pose'
```





# Quiz 2


- Another pose image is also warrior II, why it shows “unknown pose”?
- Please show the correct result to TA





# Discussion 3

- If your flexibility isn't good, how can you adjust the code to improve detection accuracy?





# Summary

- Practice Lab: 嵌入式模型 (mediapipe, video, audio, text)
- Write down the answer for discussion 1-3
  - D1: How to use average word embedding model to classify the text?
  - D2: How can we find the full list of audio events?
  - D3: How to adjust code for better detection with limited flexibility?
- Quiz:
  - Q1: Integrate Microphone Input, Speech-to-Text, and Text Classification
  - Q2: Correct the code to capture pose